Yangyi Zhang
Department of Computer Science, Georgetown University
Washington, DC 20057, USA
yangyi.zhang@georgetown.edu

2 October 2026

Dear Editors of *Machine Learning*,

I am submitting the manuscript "Neither the Judge nor the Rules: Instrument-Validity Failure in Safety-Response Classification" for consideration as a Research article.

Safety evaluations of large language models increasingly rely on automated labelling of responses, and the two instruments in general use are the LLM-as-judge and hand-authored rule-based classifiers. The critique of the first is well known and the second was adopted as its remedy. This paper audits both against an independent hand-labelled benchmark and finds that neither is a valid instrument: the LLM judge recalls none of the genuinely unsafe responses, and the rule-based classifier that was treated as ground truth is accurate on 53% of audited items and labels clean refusals as unsafe. It further shows that the agreement statistic used to document judge unreliability is not chance-corrected, and that on this corpus it reports a majority-class collapse as 62% consensus ($\kappa = 0.115$).

The paper is a measurement-validity result for machine learning evaluation. Its claim is narrow and checkable: *deterministic* does not imply *valid*, a classifier cannot validate itself, and the errors that matter in a safety instrument are the false negatives. It closes with the minimal validation protocol that would have caught each failure, including releasing the hand-labelled benchmark and both instruments' raw outputs so that the audit can be repeated by others.

I believe this fits the journal's scope in the same way that its methodological work on evaluation and reproducibility does: it identifies a widely used measurement practice as unreliable in a specific and reproducible way, quantifies the error direction, and supplies a protocol others can adopt. The paper is self-contained.

Sincerely,
Yangyi Zhang
Georgetown University