Yangyi Zhang
Department of Computer Science, Georgetown University
Washington, DC 20057, USA
yangyi.zhang@georgetown.edu

2 October 2026

Dear Editors of *Applied Intelligence*,

I am submitting the manuscript "Attribution Under a Floor: Why Single-Layer Ablation Overstates the Number of Safety-Critical Components in Fine-Tuned LLMs" for consideration as a Research article.

Single-component ablation is the standard way to localise safety degradation in fine-tuned language models: zero one attention layer or perturb one MLP neuron, and report the components whose ablation breaks safety. The paper audits an archived run of that design and shows that the reported count of safety-critical components is a floor effect rather than a mechanism. In the audited battery three of four prompts are already unsafe before any ablation, so the per-layer unsafe count saturates at three, and the reported set of "critical" layers is whichever layers flip the single remaining prompt. The neuron-level claim in the same report describes 205 safety-critical neurons at a threshold of $|\text{impact}| > 0.1$; the stored attribution results contain one neuron above that threshold, with a largest recorded magnitude of 0.156. The reported exponential decay of layer impact also does not describe the stored profile, whose largest late-layer impacts are twenty times what the fit predicts.

I believe this fits the journal's scope. It is a methodological contribution to interpretability and to applied safety auditing: it identifies a failure mode that recurs in any ablation study run on a small prompt battery, gives the corrected statistic (per-layer unsafe rate over baseline-safe prompts, with an explicit denominator), and states the battery size at which the count of critical components becomes identifiable at all. The stored artefacts are submitted alongside the manuscript so that every recomputed number can be checked.

Sincerely,
Yangyi Zhang
Georgetown University